

```
In [3]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

df = pd.read_csv("ecommerce.csv")

numeric_data = df.select_dtypes(include=['float64', 'int64'])

scaler = StandardScaler()
scaled_data = scaler.fit_transform(numeric_data)

kmeans = KMeans(n_clusters=3, random_state=42)
kmeans_labels = kmeans.fit_predict(scaled_data)

df['KMeans_Cluster'] = kmeans_labels
print(df.head())
```

	Customer ID	Gender	Age	City	Membership Type	Total Spend \
0	101	Female	29	New York	Gold	1120.20
1	102	Male	34	Los Angeles	Silver	780.50
2	103	Female	43	Chicago	Bronze	510.75
3	104	Male	30	San Francisco	Gold	1480.30
4	105	Male	27	Miami	Silver	720.40

	Items Purchased	Average Rating	Discount Applied \
0	14	4.6	True
1	11	4.1	False
2	9	3.4	True
3	19	4.7	False
4	13	4.0	True

	Days Since Last Purchase	Satisfaction Level	KMeans_Cluster
0	25	Satisfied	1
1	18	Neutral	0
2	42	Unsatisfied	0
3	12	Satisfied	1
4	55	Unsatisfied	2

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.

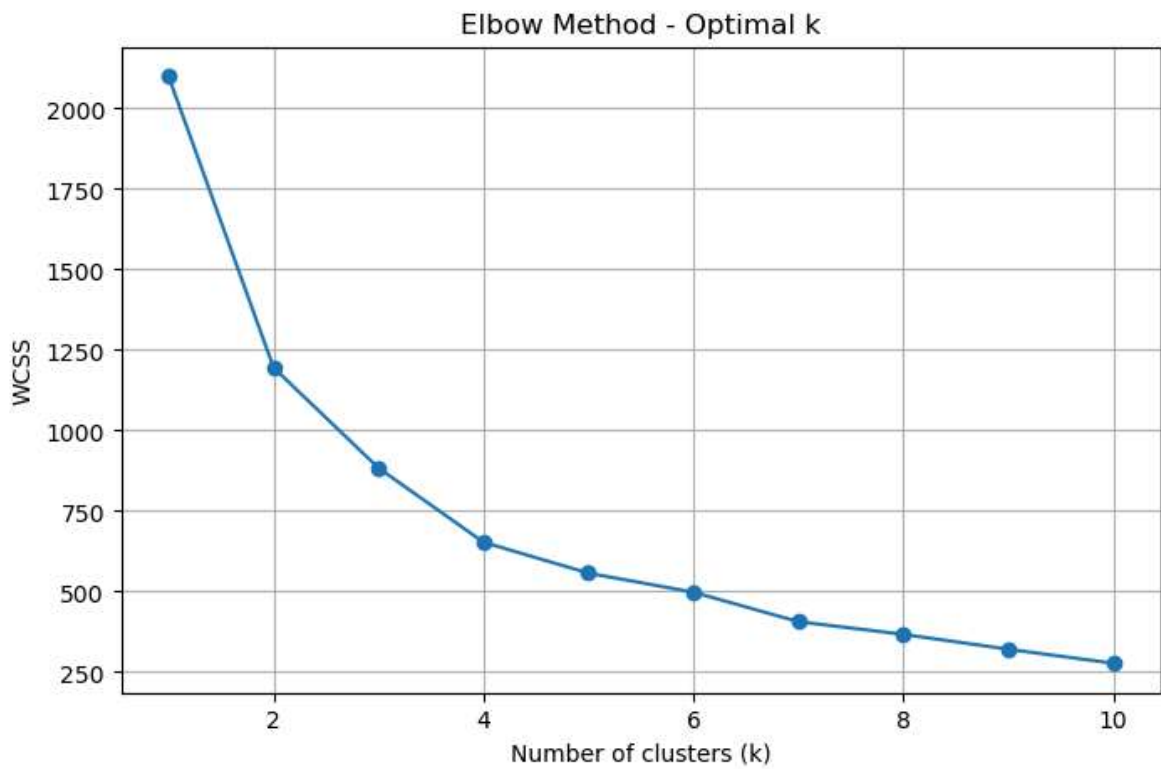
warnings.warn(

```
In [4]: import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

wcss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(scaled_data)
    wcss.append(kmeans.inertia_)

plt.figure(figsize=(8, 5))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method - Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()
```

[illegible]




```
In [5]: import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

wcss = []
K_range = range(1, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(scaled_data)
    wcss.append(kmeans.inertia_)

plt.figure(figsize=(8, 5))
plt.plot(K_range, wcss, 'bo-')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('WCSS')
plt.xticks(K_range)
plt.grid(True)
plt.show()
```

[illegible]



```
In [6]: from sklearn.cluster import AgglomerativeClustering

agglo = AgglomerativeClustering(n_clusters=3, linkage='ward')
agglo_labels = agglo.fit_predict(scaled_data)

df['Agglo_Cluster'] = agglo_labels

print(df.head())
```

	Customer ID	Gender	Age	City	Membership Type	Total Spend \
0	101	Female	29	New York	Gold	1120.20
1	102	Male	34	Los Angeles	Silver	780.50
2	103	Female	43	Chicago	Bronze	510.75
3	104	Male	30	San Francisco	Gold	1480.30
4	105	Male	27	Miami	Silver	720.40

	Items Purchased	Average Rating	Discount Applied \
0	14	4.6	True
1	11	4.1	False
2	9	3.4	True
3	19	4.7	False
4	13	4.0	True

	Days Since Last Purchase	Satisfaction Level	KMeans_Cluster	Agglo_Cluster
0	25	Satisfied	1	0
1	18	Neutral	0	0
2	42	Unsatisfied	0	1
3	12	Satisfied	1	2
4	55	Unsatisfied	2	0

```
In [10]: from sklearn.cluster import AgglomerativeClustering

agglo = AgglomerativeClustering(n_clusters=3, metric='euclidean', linkage='ward')
agglo_labels = agglo.fit_predict(scaled_data)

df['Agglo_Cluster'] = agglo_labels
```



```
print(df.head())
```

	Customer ID	Gender	Age	City	Membership Type	Total Spend \
0	101	Female	29	New York	Gold	1120.20
1	102	Male	34	Los Angeles	Silver	780.50
2	103	Female	43	Chicago	Bronze	510.75
3	104	Male	30	San Francisco	Gold	1480.30
4	105	Male	27	Miami	Silver	720.40


	Items Purchased	Average Rating	Discount Applied \
0	14	4.6	True
1	11	4.1	False
2	9	3.4	True
3	19	4.7	False
4	13	4.0	True

	Days Since Last Purchase	Satisfaction Level	KMeans_Cluster	Agglo_Cluster
0	25	Satisfied	1	0
1	18	Neutral	0	0
2	42	Unsatisfied	0	1
3	12	Satisfied	1	2
4	55	Unsatisfied	2	0

```
In [11]: import matplotlib.pyplot as plt
import scipy.cluster.hierarchy as sch

linkage_matrix = sch.linkage(scaled_data, method='ward')

plt.figure(figsize=(12, 6))
sch.dendrogram(linkage_matrix)
plt.title("Hierarchical Clustering Dendrogram")
plt.xlabel("Sample Index")
plt.ylabel("Distance")
plt.grid(True)
plt.show()
```



```
In [12]: print("K-Means Cluster Counts:")
print(df['KMeans_Cluster'].value_counts())
```

```
print("\nAgglomerative Cluster Counts:")
print(df['Agglo_Cluster'].value_counts())
```

K-Means Cluster Counts:

KMeans_Cluster

0 190

1 126

2 34

Name: count, dtype: int64

Agglomerative Cluster Counts:

Agglo_Cluster

0 176

1 116

2 58

Name: count, dtype: int64

```
In [13]: comparison = pd.crosstab(df['KMeans_Cluster'], df['Agglo_Cluster'])
print("\nCross-tabulation of KMeans vs Agglomerative Clusters:")
print(comparison)
```

Cross-tabulation of KMeans vs Agglomerative Clusters:

Agglo_Cluster	0	1	2
KMeans_Cluster			
0	74	116	0
1	68	0	58
2	34	0	0

```
In [14]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

pca = PCA(n_components=2)
pca_data = pca.fit_transform(scaled_data)

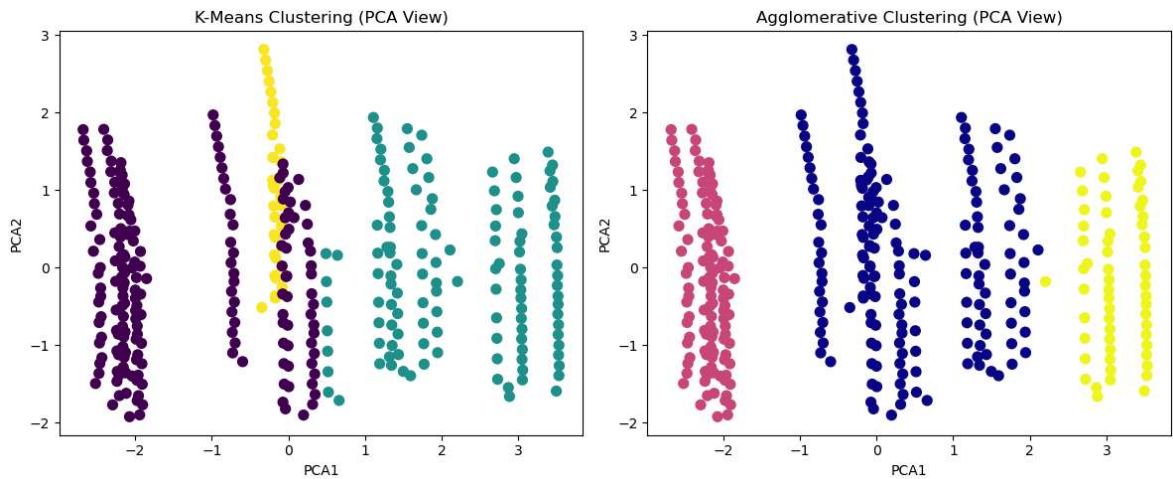
df['PCA1'] = pca_data[:, 0]
df['PCA2'] = pca_data[:, 1]

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.scatter(df['PCA1'], df['PCA2'], c=df['KMeans_Cluster'], cmap='viridis', s=50)
plt.title("K-Means Clustering (PCA View)")
plt.xlabel("PCA1")
plt.ylabel("PCA2")

plt.subplot(1, 2, 2)
plt.scatter(df['PCA1'], df['PCA2'], c=df['Agglo_Cluster'], cmap='plasma', s=50)
plt.title("Agglomerative Clustering (PCA View)")
plt.xlabel("PCA1")
plt.ylabel("PCA2")

plt.tight_layout()
plt.show()
```



In [15]: `from sklearn.decomposition import PCA`

```
pca = PCA(n_components=2)
pca_data = pca.fit_transform(scaled_data)

print("Explained variance ratio:", pca.explained_variance_ratio_)
```

Explained variance ratio: [0.61865549 0.17541405]

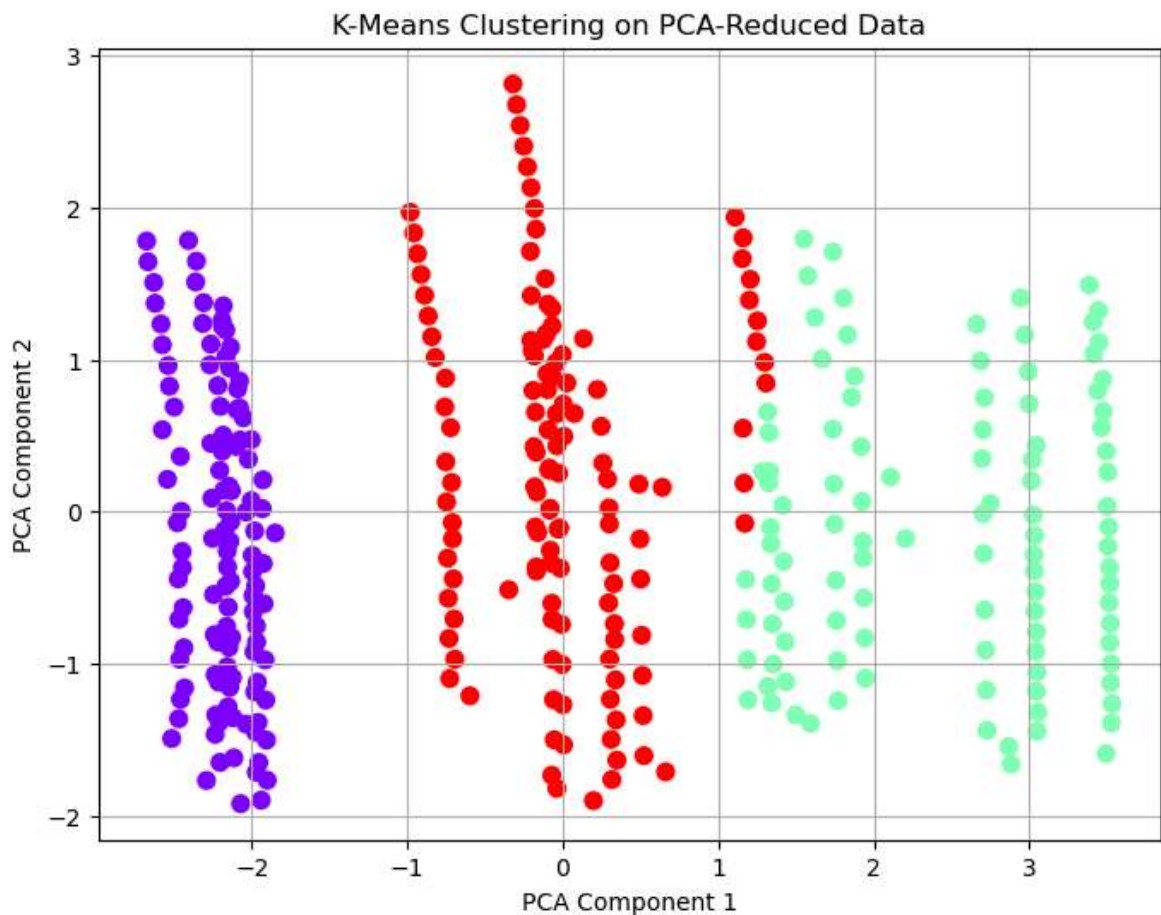
```
In [17]: kmeans_pca = KMeans(n_clusters=3, random_state=42)
kmeans_pca_labels = kmeans_pca.fit_predict(pca_data)

df['KMeans_PCA_Cluster'] = kmeans_pca_labels
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6))
plt.scatter(pca_data[:, 0], pca_data[:, 1], c=kmeans_pca_labels, cmap='rainbow',
plt.title("K-Means Clustering on PCA-Reduced Data")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.grid(True)
plt.show()
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.

```
warnings.warn(
```

```
In [18]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

pca_2d = PCA(n_components=2)
pca_data_2d = pca_2d.fit_transform(scaled_data)

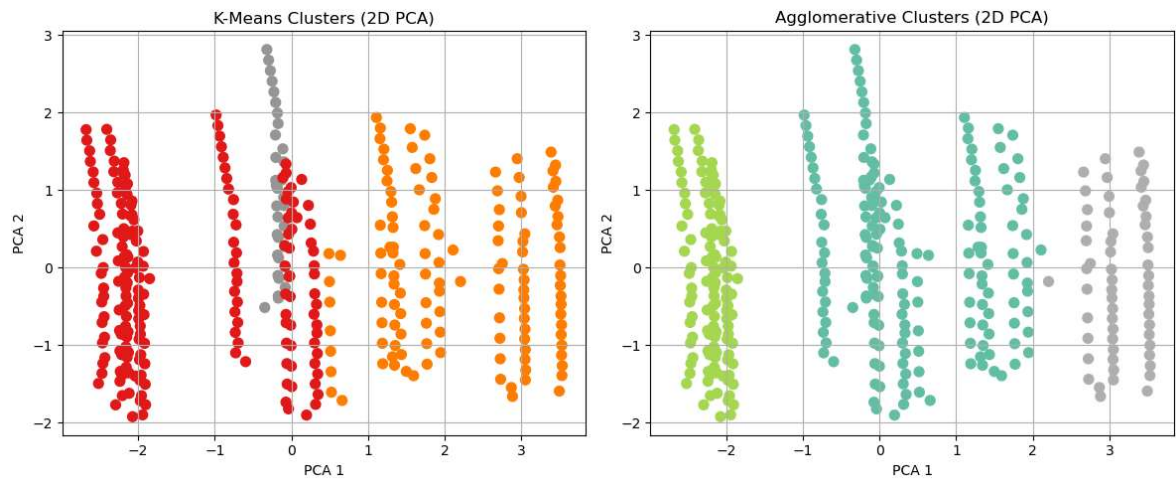
df['PCA1'] = pca_data_2d[:, 0]
df['PCA2'] = pca_data_2d[:, 1]

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.scatter(df['PCA1'], df['PCA2'], c=df['KMeans_Cluster'], cmap='Set1', s=50)
plt.title("K-Means Clusters (2D PCA)")
plt.xlabel("PCA 1")
plt.ylabel("PCA 2")
plt.grid(True)

plt.subplot(1, 2, 2)
plt.scatter(df['PCA1'], df['PCA2'], c=df['Agglo_Cluster'], cmap='Set2', s=50)
plt.title("Agglomerative Clusters (2D PCA)")
plt.xlabel("PCA 1")
plt.ylabel("PCA 2")
plt.grid(True)

plt.tight_layout()
plt.show()
```



```
In [19]: from mpl_toolkits.mplot3d import Axes3D


# Step 1: Apply PCA (3 components)
pca_3d = PCA(n_components=3)
pca_data_3d = pca_3d.fit_transform(scaled_data)

# Step 2: 3D Scatter Plot for K-Means
fig = plt.figure(figsize=(10, 6))
ax = fig.add_subplot(111, projection='3d')

scatter = ax.scatter(
    pca_data_3d[:, 0], pca_data_3d[:, 1], pca_data_3d[:, 2],
    c=df['KMeans_Cluster'], cmap='viridis', s=50
)

ax.set_title("K-Means Clusters (3D PCA)")
ax.set_xlabel("PCA 1")
ax.set_ylabel("PCA 2")
ax.set_zlabel("PCA 3")
plt.show()
```

K-Means Clusters (3D PCA)



```
In [20]: from sklearn.metrics import silhouette_score

score_kmeans = silhouette_score(scaled_data, df['KMeans_Cluster'])
print(f"Silhouette Score (K-Means): {score_kmeans:.3f}")

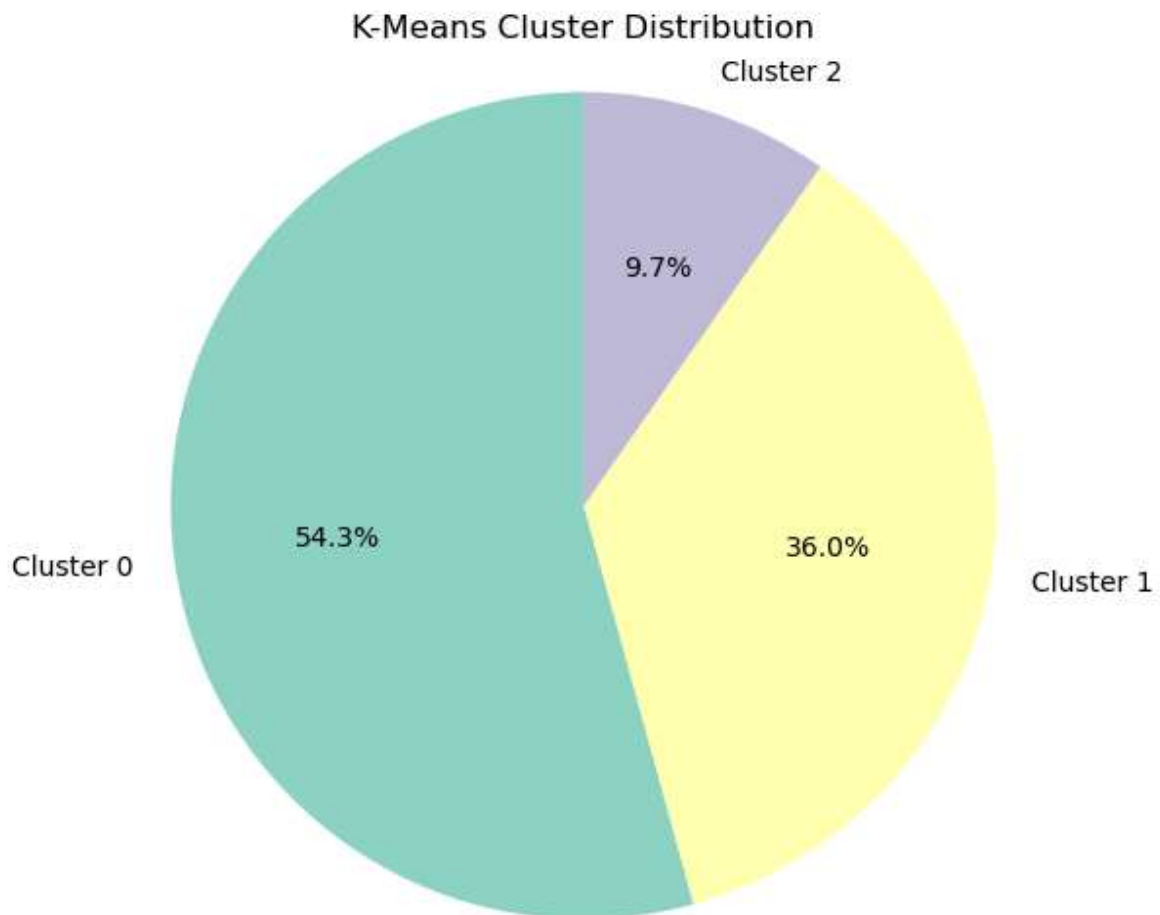
score_agglo = silhouette_score(scaled_data, df['Agglo_Cluster'])
print(f"Silhouette Score (Agglomerative): {score_agglo:.3f}")
```

Silhouette Score (K-Means): 0.392
 Silhouette Score (Agglomerative): 0.336

```
In [21]: import matplotlib.pyplot as plt

cluster_counts = df['KMeans_Cluster'].value_counts().sort_index()
labels = [f'Cluster {i}' for i in cluster_counts.index]

plt.figure(figsize=(6, 6))
plt.pie(cluster_counts, labels=labels, autopct='%1.1f%%', startangle=90, colors=
plt.title("K-Means Cluster Distribution")
plt.axis('equal')
plt.show()
```



```
In [26]: kmeans_profile = df.groupby('KMeans_Cluster').mean(numeric_only=True)

print("KMeans Cluster Profile:")
print(kmeans_profile)
```

KMeans Cluster Profile:

	Customer ID	Age	Total Spend	Items Purchased \
KMeans_Cluster				
0	279.926316	37.136842	584.933684	9.473684
1	270.992063	30.095238	1276.355556	17.269841
2	267.470588	26.794118	703.688235	12.764706

	Average Rating	Discount Applied	Days Since Last Purchase \
KMeans_Cluster			
0	3.605789	0.431579	27.736842
1	4.642857	0.468254	17.682540
2	4.017647	1.000000	53.176471

	Agglo_Cluster	PCA1	PCA2	KMeans_PCA_Cluster
KMeans_Cluster				
0	0.610526	-1.416555	-0.141727	0.778947
1	0.920635	2.183864	-0.052756	1.166667
2	0.000000	-0.177100	0.987512	2.000000

```
In [31]: kmeans_profile = df.groupby('KMeans_Cluster').mean(numeric_only=True)

print(kmeans_profile.round(2))
```

	Customer ID	Age	Total Spend	Items Purchased	\
KMeans_Cluster					
0	279.93	37.14	584.93	9.47	
1	270.99	30.10	1276.36	17.27	
2	267.47	26.79	703.69	12.76	

	Average Rating	Discount Applied	Days Since Last Purchase	\
KMeans_Cluster				
0	3.61	0.43	27.74	
1	4.64	0.47	17.68	
2	4.02	1.00	53.18	

	Agglo_Cluster	PCA1	PCA2	KMeans_PCA_Cluster
KMeans_Cluster				
0	0.61	-1.42	-0.14	0.78
1	0.92	2.18	-0.05	1.17
2	0.00	-0.18	0.99	2.00

```
In [32]: cluster_labels = {
          0: 'High Spenders',
          1: 'Lost',
          2: 'Loyal'
        }

df['Customer_Segment'] = df['KMeans_Cluster'].map(cluster_labels)

df[['KMeans_Cluster', 'Customer_Segment']].head()
```

```
Out[32]:
```

	KMeans_Cluster	Customer_Segment
0	1	Lost
1	0	High Spenders
2	0	High Spenders
3	1	Lost
4	2	Loyal

```
In [33]: print(df['Customer_Segment'].value_counts())
```

```
Customer_Segment
High Spenders    190
Lost             126
Loyal            34
Name: count, dtype: int64
```

```
In [ ]:
```